

April 9, 2026

Andre van Zijl

Andre,

I read your GPU On-Ramp paper with real interest. You are defining a category that we have been building product against for the past two years, and I wanted to introduce MaiaEdge and show you where our work overlaps.

MaiaEdge in 30 Seconds

MaiaEdge is carrier infrastructure for federated private networking. We make a 1RU edge appliance (the Path Border Controller) and a cloud-based orchestrator (the Path Computation Engine) that together let any network operator provision deterministic, encrypted, private Ethernet paths across network boundaries in minutes, not months. No BGP, no MPLS, no manual coordination. Our founding team built Acme Packet (acquired by Oracle, \$2.1B) and 128 Technology (acquired by Juniper Networks, \$450M). This is our third company solving carrier infrastructure problems.

Where Your Thesis Meets Our Product

Your paper identifies four layers that define GPU On-Ramp capability: power, cooling, interconnection, and flexibility. MaiaEdge is purpose-built to make the interconnection layer operational. Specifically:

- **The automation gap you leave unnamed.** You describe carrier-neutral fabric with 100G-400G cross-connects and neocloud on-ramps. The reality underneath that is manual NNI builds, VLAN spreadsheets, and 60-90 day provisioning cycles. MaiaEdge automates that entire workflow, turning GPU On-Ramp interconnection from a facility attribute into an operational capability that scales.
- **Inference demands deterministic paths.** Your Section 2.2 argues inference workloads cannot tolerate the latency of distant training campuses. We agree, and go further: they cannot tolerate the jitter of best-effort routing either. MaiaEdge delivers deterministic, SLA-aware paths with sub-2-microsecond forwarding overhead and line-rate AES-256 encryption. The network becomes part of the compute loop.

- **Neocloud distribution at speed.** You identify the neocloud paradox: demand growing 200%+ but infrastructure taking years to build. MaiaEdge solves the connectivity side of that equation. A neocloud deploying into a new colo facility can light up private, multi-tenant enterprise on-ramps through MaiaEdge in minutes rather than negotiating individual NNIs with every carrier.
- **Multi-cloud hybrid architecture.** Your enterprise SaaS scenario (Section 6.4) describes hyperscaler and neocloud environments interconnected inside the same facility. MaiaEdge does this today via native API integrations with Equinix Fabric and Megaport, giving operators a single control plane for cloud on-ramp delivery under their own brand.
- **Sovereignty without surrender.** Your enterprise use cases (finance, healthcare, SaaS) all require private interconnection where regulated data never touches the public internet. MaiaEdge's architecture ensures the operator, not a third-party fabric, owns the customer relationship, the pricing, and the data path.

Why I Am Reaching Out

Your framework gives the industry a shared vocabulary for something we see every day in our sales conversations with colocation operators and neoclouds. We think there is a conversation worth having about how MaiaEdge fits into the GPU On-Ramp stack as the interconnection automation layer.

My co-founder Abilash Menon (our CEO) and I will be in Dallas next week. If you are in the area, I would welcome the chance to connect in person. If not, happy to set up a call at your convenience.

Best,

Timothy Ziemer

CRO & Co-Founder, MaiaEdge

timziemer@maiaedge.io

maiaedge.io